

# Parameter estimation Optimization

## Lecture 2d

# Parametric functions

- General ML model  $Y \sim \text{Distribution}(f_{\theta}(\mathbf{x}), \dots)$ 
  - Some of them:  $Y \sim f_{\theta}(\mathbf{x}) + \epsilon, \epsilon \sim \text{Distribution}(\dots)$
  - Generalization of simple models can be done
- **Example:** logistic  $Y \sim \text{Bernoulli}\left(\frac{1}{1+e^{-\theta^T \mathbf{x}}}\right)$ 
  - Generalization 1: (**basis function expansion**):  
 $Y \sim \text{Bernoulli}\left(\frac{1}{1+e^{-\theta^T \phi(\mathbf{x})}}\right)$
  - Generalization 2:  $Y \sim \text{Bernoulli}\left(\frac{1}{1+e^{-f_{\theta}(\mathbf{x})}}\right)$ 
    - $f_{\theta}(\mathbf{x}) = \|\mathbf{x} - \boldsymbol{\theta}\|^2$

# Loss minimization

- Given training set  $T$ , we **want** to minimize

$$E_{new} = \int_{(\mathbf{x}_*, y_*)} E(y_*, \hat{y}(\mathbf{x}_*, T, \boldsymbol{\theta})) p(\mathbf{x}_*, y_*) d\mathbf{x}_* dy_*$$

- We **can** minimize cost function

$$J(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n L(y_i, \hat{y}(\mathbf{x}_i, \boldsymbol{\theta}))$$

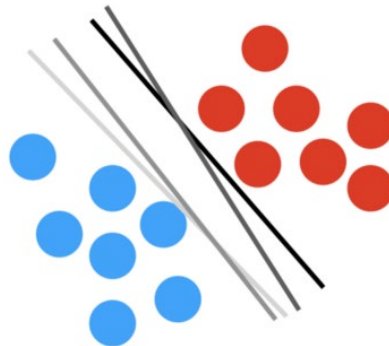
$$E_{new} \approx J(\boldsymbol{\theta})?$$

- Optimizing  $J(\boldsymbol{\theta})$  does not lead to optimizing  $E_{new}$ 
  - Overfitting



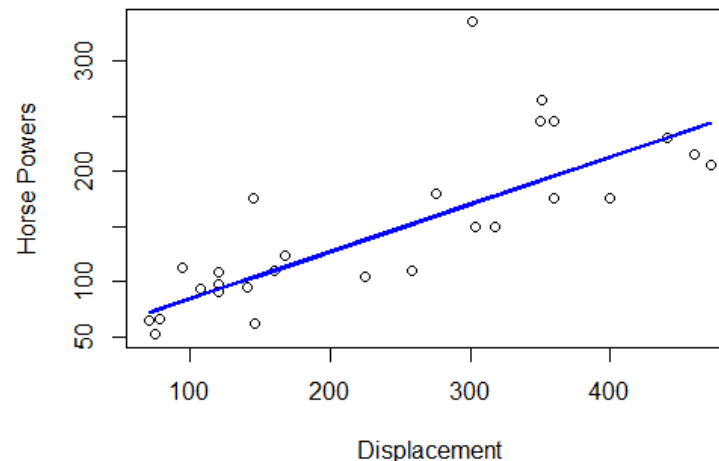
# Loss minimization: comments

- Training a model with perfect accuracy unreasonable
  - Statistical noise for finite  $n$
- Loss function can be different from error function
- Some loss functions are not good for training, for ex. misclassification rate.



# Loss functions

- Assuming a distribution, **derive as minus log-likelihood**:
- $y \sim \text{Normal}(f_\theta(x), \sigma^2) \rightarrow L(y, f_\theta(x)) = (y - f_\theta(x))^2$
- Heavy outliers  $y \sim \text{Laplace}(f_\theta(x), r) \rightarrow L(y, f_\theta(x)) = |y - f_\theta(x)|$

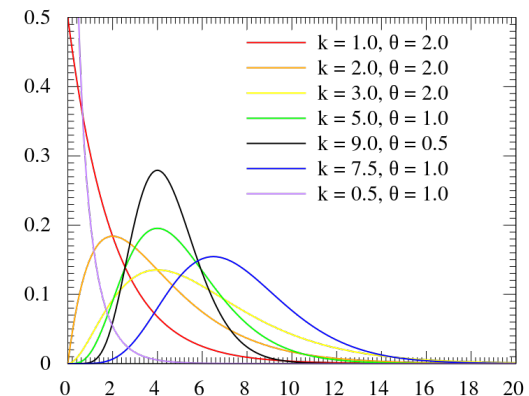
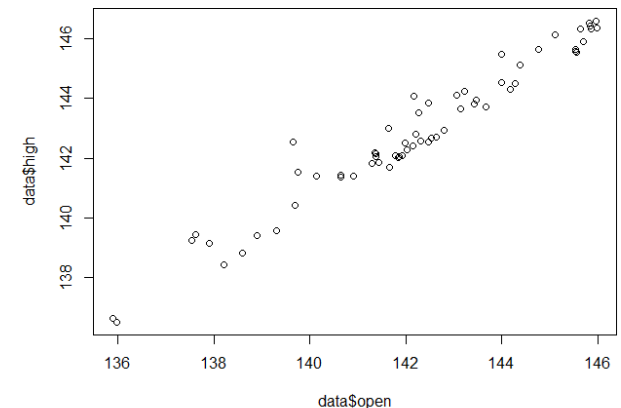
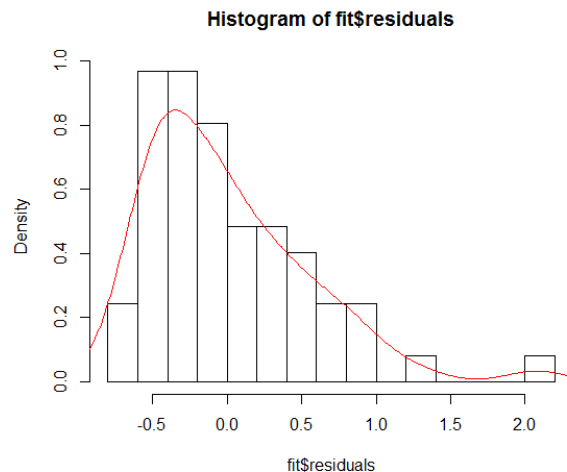


# Loss functions

- Count data  $y \sim \text{Poisson}(f_\theta(x))$

**Example:** Daily Stock prices NASDAQ

- Open
  - High (within day)
1. Try to fit usual linear regression, study histogram of residuals



# Loss functions

- If the distribution is difficult to assume / only some properties known → **ad-hoc loss functions**
- **Huber loss**: similar to quadratic but robust to outliers

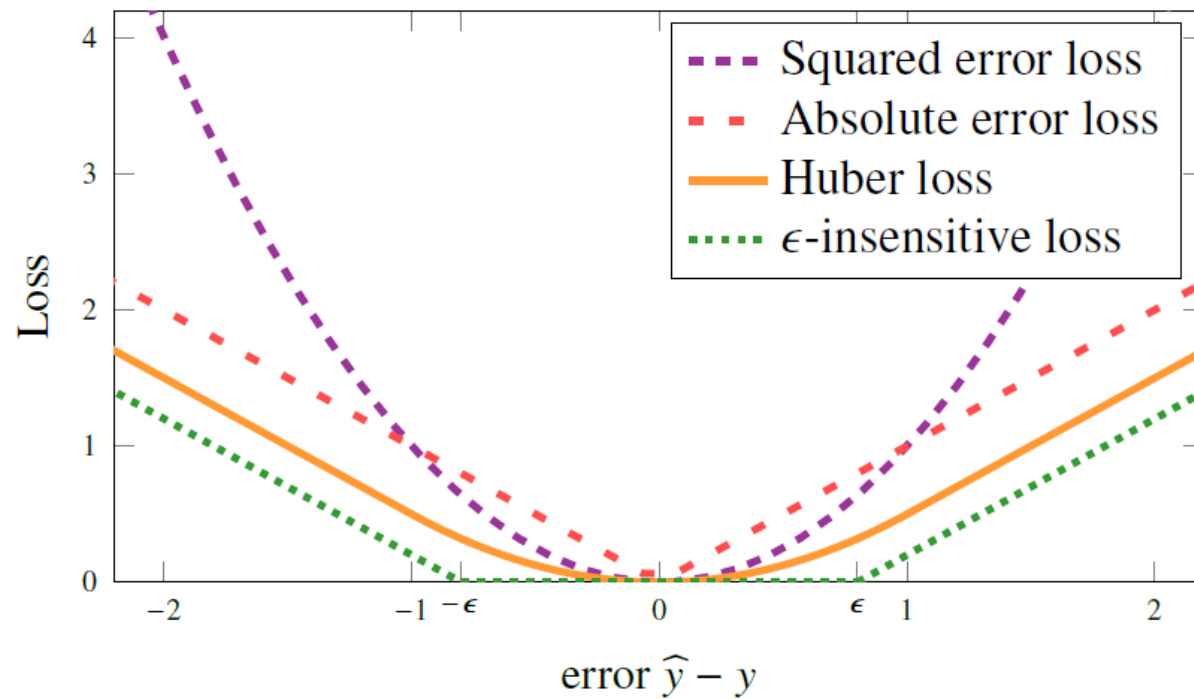
$$L(y, \hat{y}) = \begin{cases} \frac{1}{2}(\hat{y} - y)^2 & \text{if } |\hat{y} - y| < 1, \\ |\hat{y} - y| - \frac{1}{2} & \text{otherwise.} \end{cases}$$

- **E-intensive loss**

$$L(y, \hat{y}) = \begin{cases} 0 & \text{if } |\hat{y} - y| < \epsilon, \\ |\hat{y} - y| - \epsilon & \text{otherwise,} \end{cases}$$



# Loss functions





# Loss functions: classification


- **Cross-entropy** corresponds to minus log-likelihood:

$$J(y, \hat{p}(y)) = - \sum_{i=1}^n \sum_{m=1}^M I(y_i = C_m) \log \hat{p}(y_i = C_m)$$

- Ad-hoc loss functions binary classification  $\mathcal{C} = \{-1, 1\}$ 
  - Assume model returns  $f(\mathbf{x})$ :  $\hat{y} = \text{sign}(f(\mathbf{x}))$

- **Example:** logistic  $f(\mathbf{x}) = \frac{1}{1+e^{-\theta^T \mathbf{x}}} - 0.5$

- **Note:** mistake when  $yf(\mathbf{x}) = -1$



# Loss functions: classification

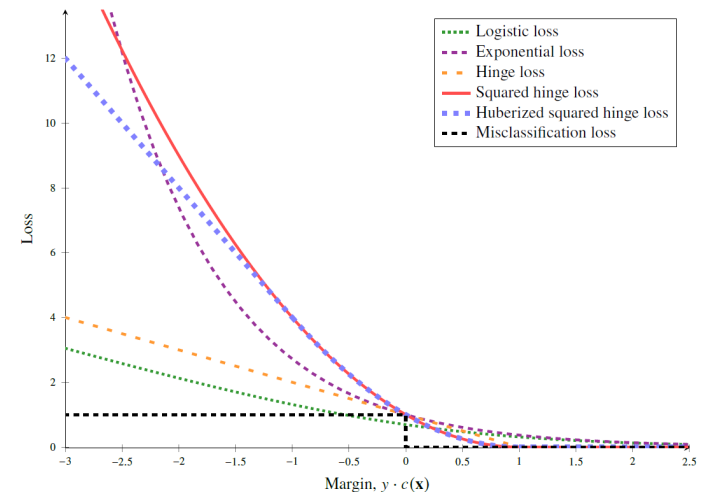
## Ad-hoc loss functions binary classification

- Exponential loss**

$$L(y \cdot f(\mathbf{x})) = \exp(-y \cdot f(\mathbf{x}))$$

- Hinge loss**

$$L(y \cdot f(\mathbf{x})) = \begin{cases} 1 - y \cdot f(\mathbf{x}) & \text{for } y \cdot f(\mathbf{x}) \leq 1 \\ 0 & \text{otherwise.} \end{cases}$$



# Loss functions: classification

## Binary to multiclass

- **One versus one**: class  $C_i$  vs class  $C_j$  + majority voting from all classifiers
- **One versus rest**: class  $C_i$  vs not  $C_i$  + highest probability class
- **Comparison**: OVO needs less data to train one model but more models.

# Regularization

- $E_{new} \approx J(\theta)$ ? – no
- Similar for (moderately) simple models, not similar for too complex model (overfitting).
- **Explicit regularization**: penalize complexity by changing cost function
- **Implicit regularization**: **early stopping**
  - If cost function optimized iteratively, don't let it decrease too much



# Explicit regularization

- Penalize cost function

$$\min_{\theta} J(\theta) + \lambda R(\theta)$$

- $\lambda > 0$

- **L1 regularization:**  $R(\theta) = \lambda \|\theta\|_1$
- **L2 regularization:**  $R(\theta) = \lambda \|\theta\|_2$

- **Example:** Ridge regression

$$\min_{\theta} \frac{1}{n} \sum_{i=1}^n (\mathbf{y}_i - \theta^T \mathbf{x}_i)^2 + \lambda \sum_{j=1}^p \theta_j^2, \quad \lambda > 0$$

# Explicit regularization: ridge regression

Equivalent form

$$\hat{\theta}^{ridge} = \operatorname{argmin} \sum_{i=1}^N (y_i - \theta_0 - \theta_1 x_{1i} - \dots - \theta_p x_{pi})^2$$

subject to  $\sum_{j=1}^p \theta_j^2 \leq s$

*Solution*

$$\theta^{ridge} = (X^T X + \lambda I)^{-1} X^T y$$

# Ridge regression

## Properties

- Extreme cases:
  - $\lambda = 0$  usual linear regression (no shrinkage)
  - $\lambda = +\infty$  fitting a constant ( $\theta = 0$  except of  $\theta_0$ )
- Degrees of freedom decrease when  $\lambda$  increases
  - $\lambda = 0 \rightarrow d.f. = p$
- $p > n$  is doable
  - Compare with linear regression
- How to estimate  $\lambda$ ?
  - cross-validation

# Ridge regression

**Example** Computer Hardware Data Set : performance measured for various processors and also

- Cycle time
- Memory
- Channels
- ...

Build model predicting performance






# Ridge regression

- R code: use package **glmnet** with  $\alpha=0$  (Ridge regression)
- Seeing how Ridge converges

```
data=read.csv("machine.csv", header=F)
```

```
library(caret)  
library(glmnet)  
scaler=preProcess(data)  
data1=predict(scaler, data)  
covariates=data1[,3:8]  
response=data1[, 9]
```

```
model0=glmnet(as.matrix(covariates),  
response, alpha=0,family="gaussian")  
plot(model0, xvar="lambda", label=TRUE)
```

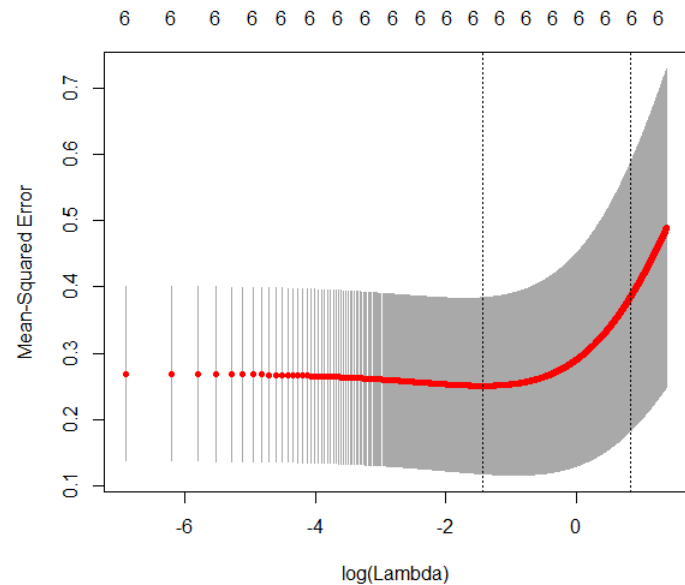


# Ridge regression

- Choosing the best model by cross-validation:

```
model=cv.glmnet(as.matrix(covariates),  
response, alpha=0,family="gaussian")  
model$lambda.min  
plot(model)  
coef(model, s="lambda.min")
```

```
> coef(model, s="lambda.min")  
7 x 1 sparse Matrix of class "dgCMatrix"  
1  
(Intercept) -4.530442e-17  
v3          3.420739e-02  
v4          3.085696e-01  
v5          3.403839e-01  
v6          1.593470e-01  
v7          5.489116e-02  
v8          1.970982e-01
```



```
> model$lambda.min  
[1] 0.046
```

# Ridge regression

- How good is this model in prediction?

```
covariates=train[,1:6]
response=train[, 7]
model=cv.glmnet(as.matrix(covariates), response, alpha=1,family="gaussian",
lambda=seq(0,1,0.001))
y=test[,7]
ynew=predict(model, newx=as.matrix(test[, 1:6]), type="response")
```

```
#Coefficient of determination
sum((ynew-mean(y))^2)/sum((y-mean(y))^2)
```

```
sum((ynew-y)^2)
```

Note that data are so small so numbers  
change much for other train/test

```
> sum((ynew-mean(y))^2)/sum((y-mean(y))^2)
[1] 0.5438148
> sum((ynew-y)^2)
[1] 18.04988
> |
```



# LASSO

- Add  $l_1$  regularization term

$$\hat{\theta}^{lasso} = \operatorname{argmin} \left\{ \frac{1}{n} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_{1j} - \dots - \theta_p x_{pj})^2 + \lambda \sum_{j=1}^p |\theta_j| \right\}$$

- $\lambda > 0$  is **penalty factor**
- Equivalent formulation

$$\begin{aligned} \hat{\theta}^{lasso} &= \operatorname{argmin} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_{1j} - \dots - \theta_p x_{pj})^2 \\ \text{subject to } &\sum_{j=1}^p |\theta_j| \leq s \end{aligned}$$

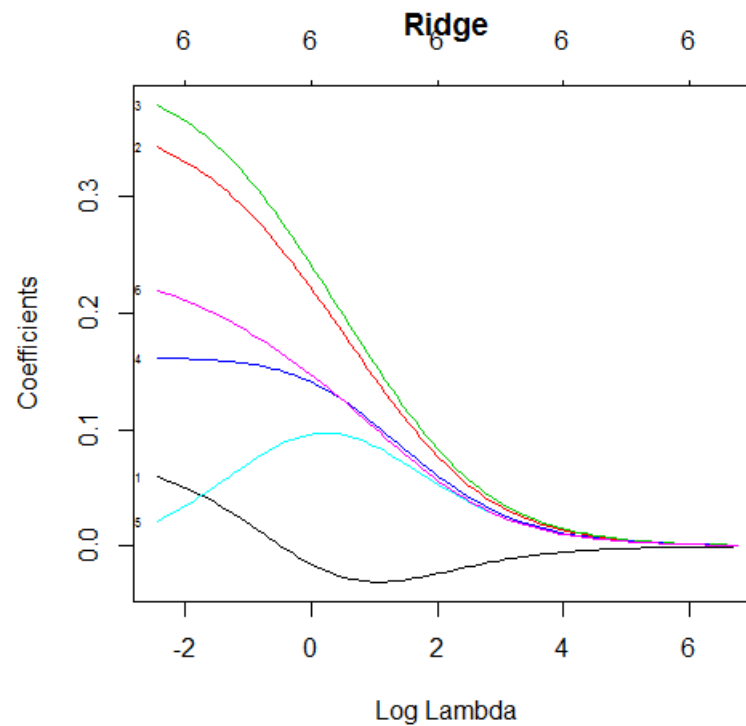
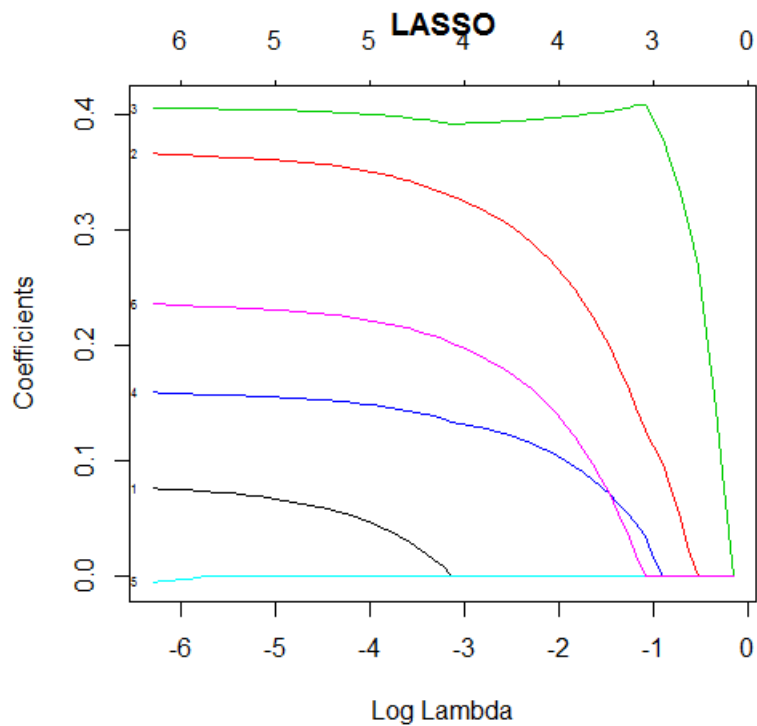




# LASSO vs Ridge

- **LASSO yields sparse solutions!**

**Example** Computer hardware data



# LASSO vs Ridge

- In R, use glmnet with **alpha=1**
- Only 5 variables selected by LASSO

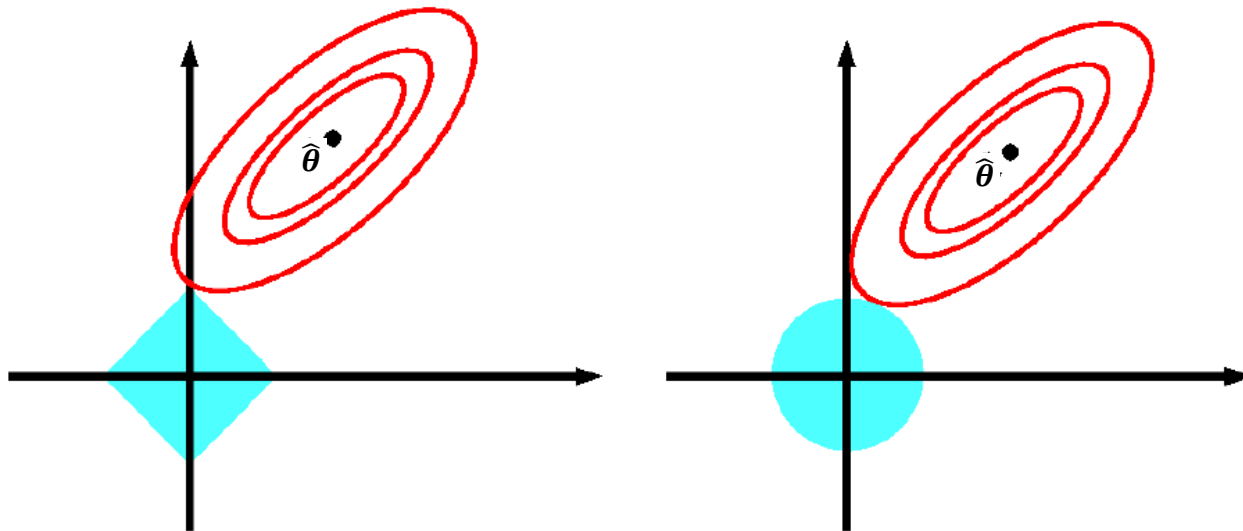
```
> coef(model, s="lambda.min")  
7 x 1 sparse Matrix of class "dgCMatrix"
```

```
      1  
(Intercept) -5.091825e-17  
v3           6.350488e-02  
v4           3.578607e-01  
v5           4.033670e-01  
v6           1.541329e-01  
v7           .  
v8           2.287134e-01  
> |
```

```
> sum((ynew-mean(y))^2)/sum((y-mean(y))^2)  
[1] 0.5826904  
> sum((ynew-y)^2)  
[1] 16.63756
```

# LASSO vs Ridge

- Why Lasso leads to sparse solutions?
  - Feasible area for Ridge is a circle (2D)
  - Feasible area for LASSO is a polygon (2D)



# LASSO properties

- **Lasso is widely used when  $p \gg n$** 
  - Linear regression breaks down when  $p > n$
  - Application: DNA sequence analysis, Text Prediction
- No explicit formula for  $\hat{\theta}^{lasso}$ 
  - Optimization algorithms used



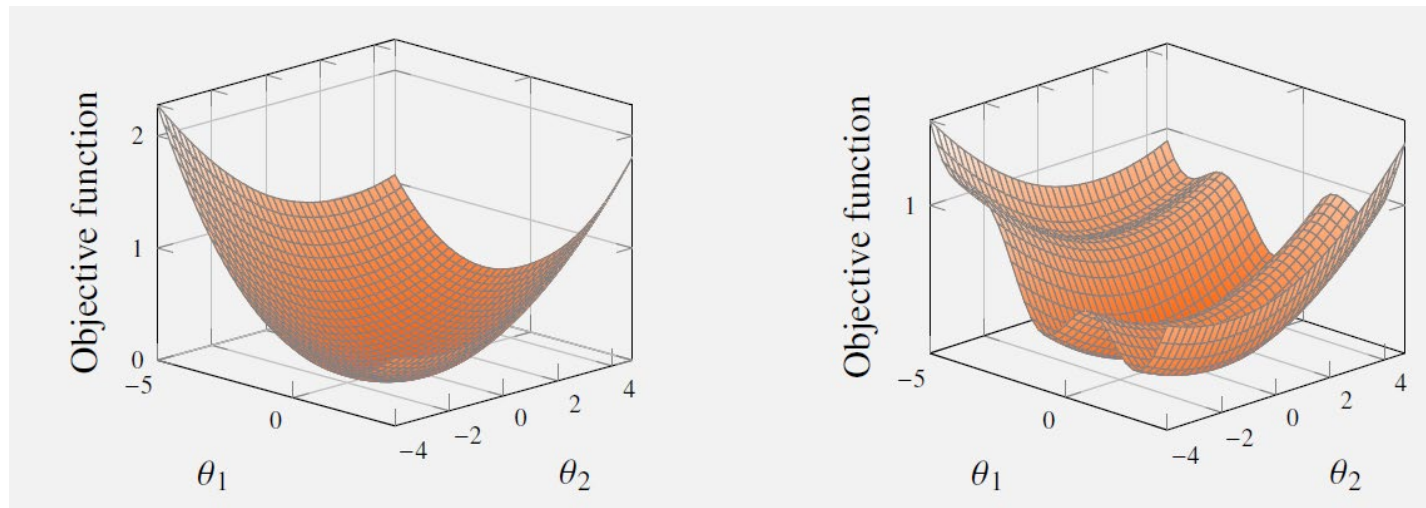
# Optimization methods

- Numerical optimization often needed

$$\min_{\theta} J(\theta)$$

$$\min_{\lambda} E_{hold-out}(\lambda)$$

- If not convex objective, more than one local optimum



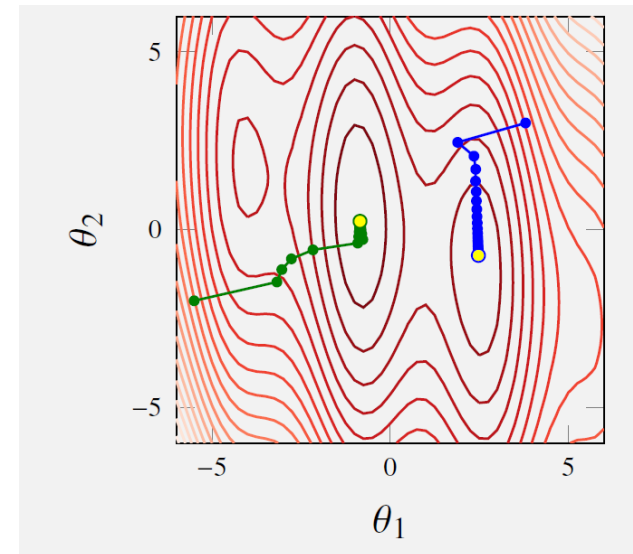
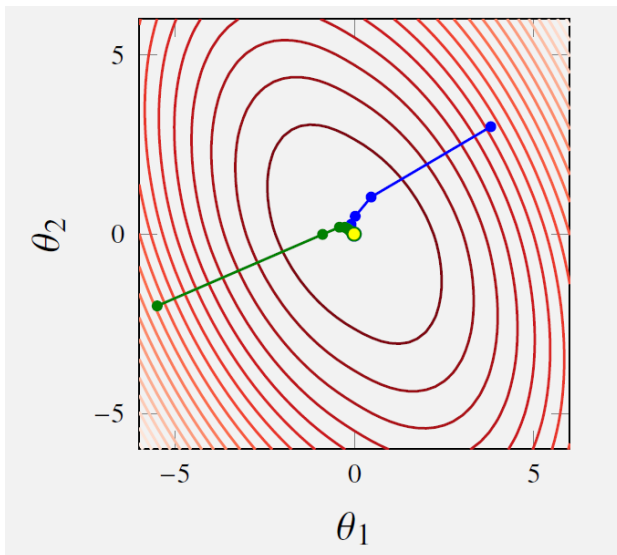
# Optimization methods

- **Gradient descent method**

$$\hat{\theta} = \arg \min_{\theta} J(\theta)$$

- Basic idea:

- Start from some point  $\theta_0$
- Move to the next point along **descent direction**  $-\nabla_{\theta} J(\theta)$



# Gradient descent

---

**Algorithm 5.1:** Gradient descent

---

**Input:** Objective function  $J(\theta)$ , initial  $\theta^{(0)}$ , learning rate  $\gamma$

**Result:**  $\hat{\theta}$

```
1 Set  $t \leftarrow 0$ 
2 while  $\|\theta^{(t)} - \theta^{(t-1)}\|$  not small enough do
3   |   Update  $\theta^{(t+1)} \leftarrow \theta^{(t)} - \gamma \nabla_{\theta} J(\theta^{(t)})$ 
4   |   Update  $t \leftarrow t + 1$ 
5 end
6 return  $\hat{\theta} \leftarrow \theta^{(t-1)}$ 
```

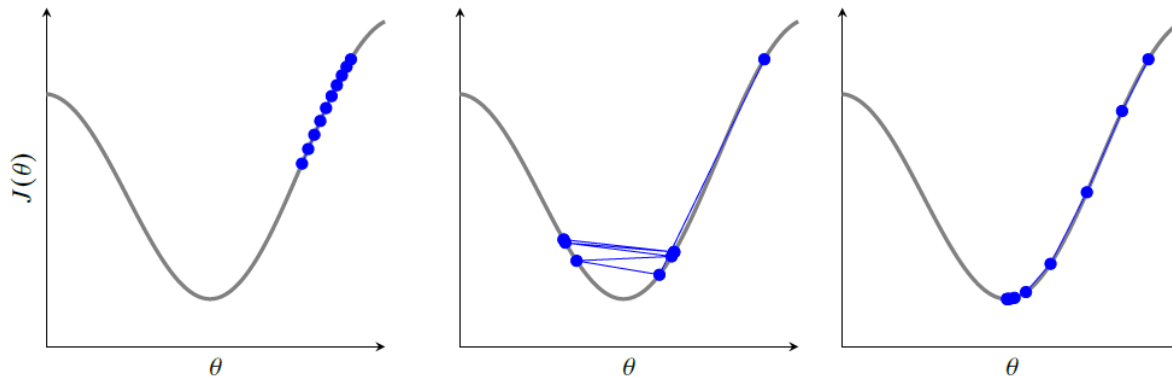
---

- **Example:** logistic regression



# Gradient descent

- Influence of  $\gamma$



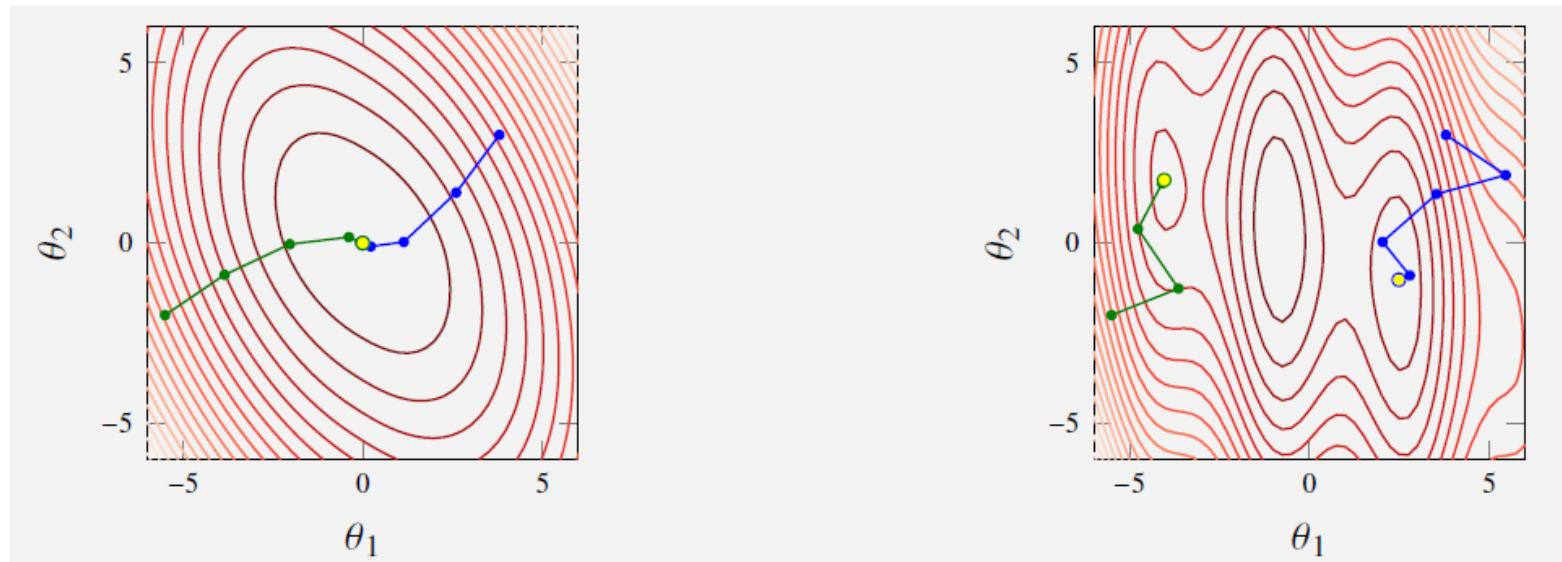
- Trace  $J(\theta^{(t)})$  vs  $t$ 
  - High oscillation  $\rightarrow$  decrease  $\gamma$
  - Slow changes  $\rightarrow$  increase  $\gamma$
- Try with different  $\theta^{(0)}$  if possible



# Newton's method

- Assume  $J(\boldsymbol{\theta})$  is "locally" quadratic
- Newton's method: move along the best direction

$$\boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \eta [\nabla_{\boldsymbol{\theta}}^2 J(\boldsymbol{\theta}^{(t)})]^{-1} [\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}^{(t)})]$$



# Newton's method

- Properties
  - No convergence guarantees
  - **Advantage**: if  $J(\boldsymbol{\theta})$  is quadratic and  $\eta = 1 \rightarrow$  convergence in one iteration
  - **Disadvantage 1**: Hessian must be invertable
  - **Disadvantage 2**: Hessian is computationally heavy
- Solution: quasi-Newton methods (ex. **BFGS**)
  - Choose some  $H^{(0)}$
  - Approximate the inverse Hessian
$$H^{(t)} = \phi(H^{(t-1)}, \nabla J(\boldsymbol{\theta}^{(t-1)}), \nabla J(\boldsymbol{\theta}^{(t)}))$$

# Newton's method

---

**Algorithm 5.2:** Trust-region Newton's method

---

**Input:** Objective function  $J(\boldsymbol{\theta})$ , initial  $\boldsymbol{\theta}^{(0)}$ , trust region radius  $D$

**Result:**  $\hat{\boldsymbol{\theta}}$

- 1 Set  $t \leftarrow 0$
  - 2 **while**  $\|\boldsymbol{\theta}^{(t)} - \boldsymbol{\theta}^{(t-1)}\|$  *not small enough* **do**
  - 3     Compute  $\mathbf{v} \leftarrow [\nabla_{\boldsymbol{\theta}}^2 J(\boldsymbol{\theta}^{(t)})]^{-1} [\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}^{(t)})]$
  - 4     Compute  $\eta \leftarrow \frac{D}{\max(\|\mathbf{v}\|, D)}$
  - 5     Update  $\boldsymbol{\theta}^{(t+1)} \leftarrow \boldsymbol{\theta}^{(t)} - \eta \mathbf{v}$
  - 6     Update  $t \leftarrow t + 1$
  - 7 **end**
  - 8 **return**  $\hat{\boldsymbol{\theta}} \leftarrow \boldsymbol{\theta}^{(t-1)}$
-

# Optimization methods in R

- In R, use `optim(par, fn, gr, method,...)`
  - `par`: initial parameter vector
  - `fn`: function to optimize
  - `gr`: gradient function
  - `method`

**Example:** trace plot for  $y = (x_1 - 2)^4 + (x_2 - 4)^4$



# Optimization methods in R

#Workaround: optim does not return iterations

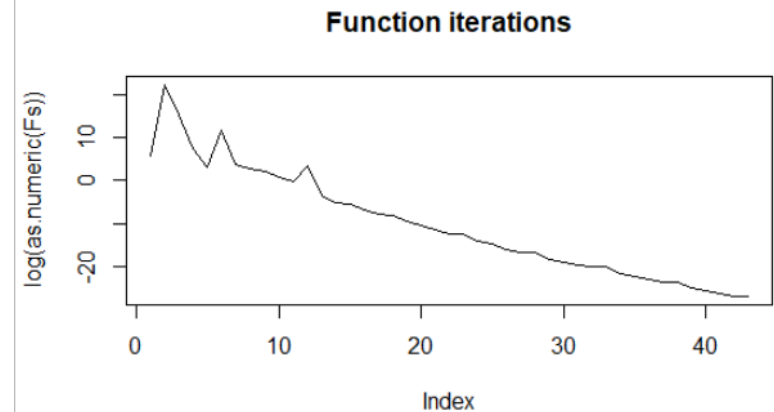
```
Fs=list()
Params=list()
k=0

myf<- function(x){
  f=(x[1]-2)^4+(x[2]-4)^4
  .GlobalEnv$k= .GlobalEnv$k+1
  .GlobalEnv$Fs[[k]]=f
  .GlobalEnv$Params[[k]]=x
  return(f)
}

myGrad <-function(x) c(4*(x[1]-2)^3, 4*(x[2]-4)^3)

res<-optim(c(0,0), fn=myf, gr=myGrad, method="BFGS")

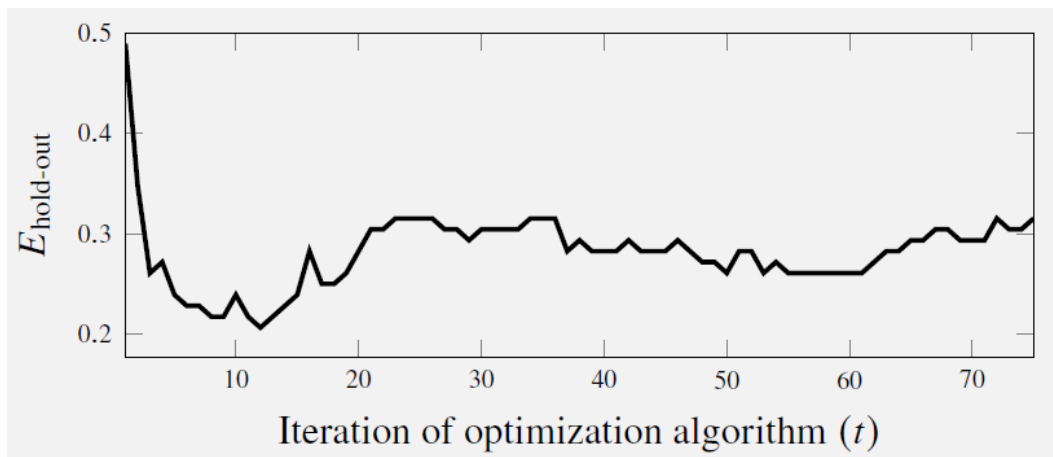
plot(log(as.numeric(Fs)), type="l", main="Function
iterations")
```



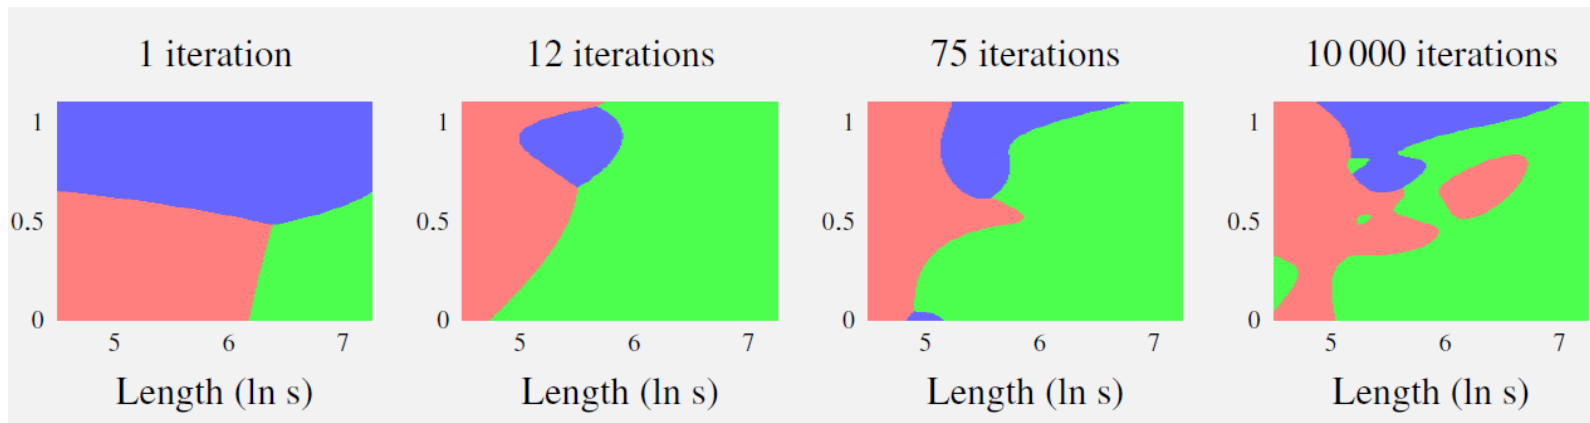
# Implicit regularization

- **Early stopping**

- For complex models, accurate model optimization may lead to overfitting
- Start from some parameter set (probably not optimal, large  $E_{train}$  and  $E_{new}$ )
- Trace the validation error (and training error? ) for each  $t$
- Choose model with the smallest validation error



# Implicit regularization



# Optimization for large data

## Stochastic gradient descent

**Idea:** use gradient descent + approximation to expected value

- For **random** sample of size  $n_b$  from sample of size  $n$

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \approx \frac{1}{n_b} \sum_{i=1}^{n_b} x_i$$
$$\nabla_{\theta} J(\theta) \approx \frac{1}{n_b} \sum_{(x_i, y_i) \in \text{sample}} \nabla_{\theta} L(x_i, y_i, \theta)$$

1. One **epoch**:
  1. Permute data and divide into batches of size  $n_b$
  2. In each optimization iteration, use one batch
2. Repeat step 1



# Stochastic gradient descent

---

**Algorithm 5.3:** Stochastic gradient descent

---

**Input:** Objective function  $J(\theta) = \frac{1}{n} \sum_{i=1}^n L(\mathbf{x}_i, y_i, \theta)$ , initial  $\theta^{(0)}$ , learning rate  $\gamma^{(t)}$

**Result:**  $\hat{\theta}$

```
1 Set  $t \leftarrow 0$ 
2 while Convergence criteria not met do
3   for  $i = 1, 2, \dots, E$  do
4     Randomly shuffle the training data  $\{\mathbf{x}_i, y_i\}_{i=1}^n$ 
5     for  $j = 1, 2, \dots, \frac{n}{n_b}$  do
6       Approximate the gradient using the mini-batch  $\{(\mathbf{x}_i, y_i)\}_{i=(j-1)n_b+1}^{jn_b}$ ,
7          $\hat{\mathbf{d}}^{(t)} = \frac{1}{n_b} \sum_{i=(j-1)n_b+1}^{jn_b} \nabla_{\theta} L(\mathbf{x}_i, y_i, \theta^{(t)})$ .
8       Update  $\theta^{(t+1)} \leftarrow \theta^{(t)} - \gamma^{(t)} \hat{\mathbf{d}}^{(t)}$ 
9       Update  $t \leftarrow t + 1$ 
10    end
11  end
12 return  $\hat{\theta} \leftarrow \theta^{(t-1)}$ 
```

---

- Different choices for  $\gamma_t$ , for ex  $\gamma^{(t)} = \frac{1}{t^\alpha}, \alpha \in (0.5, 1]$

# Hyperparameter optimization

- $E_{hold-out}$  costly to compute  $\rightarrow$  usual optimization very hard
  - Note: for each  $\lambda$  first we need to optimize  $\theta$ ...+ gradients of  $E_{hold-out}$
- Grid search (can also be costly)
  - Alternative: Bayesian optimization

